

```
In [478]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import re
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
```

```
In [608]: a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\madrid_2013.csv")
a
```

Out[608]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2013-11-01 01:00:00	NaN	0.6	NaN	NaN	135.0	74.0	NaN	NaN	NaN	7.0	NaN	NaN	28079004
1	2013-11-01 01:00:00	1.5	0.5	1.3	NaN	71.0	83.0	2.0	23.0	16.0	12.0	NaN	8.3	28079008
2	2013-11-01 01:00:00	3.9	NaN	2.8	NaN	49.0	70.0	NaN	NaN	NaN	NaN	NaN	9.0	28079011
3	2013-11-01 01:00:00	NaN	0.5	NaN	NaN	82.0	87.0	3.0	NaN	NaN	NaN	NaN	NaN	28079016
4	2013-11-01 01:00:00	NaN	NaN	NaN	NaN	242.0	111.0	2.0	NaN	NaN	12.0	NaN	NaN	28079017
...
209875	2013-03-01 00:00:00	NaN	0.4	NaN	NaN	8.0	39.0	52.0	NaN	NaN	NaN	NaN	NaN	28079056
209876	2013-03-01 00:00:00	NaN	0.4	NaN	NaN	1.0	11.0	NaN	6.0	NaN	2.0	NaN	NaN	28079057
209877	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	4.0	75.0	NaN	NaN	NaN	NaN	NaN	28079058
209878	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	11.0	52.0	NaN	NaN	NaN	NaN	NaN	28079059
209879	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	1.0	10.0	75.0	3.0	NaN	NaN	NaN	NaN	28079060

209880 rows × 14 columns

```
In [609]: a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209880 entries, 0 to 209879
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        209880 non-null  object 
 1   BEN         50462 non-null   float64
 2   CO          87018 non-null   float64
 3   EBE         50463 non-null   float64
 4   NMHC        25935 non-null   float64
 5   NO          209108 non-null   float64
 6   NO_2        209108 non-null   float64
 7   O_3         121858 non-null   float64
 8   PM10        104339 non-null   float64
 9   PM25        51980 non-null   float64
10   SO_2        86970 non-null   float64
11   TCH         25935 non-null   float64
12   TOL         50317 non-null   float64
13   station     209880 non-null   int64  
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [610]: b=a.fillna(value=55)
b
```

Out[610]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2013-11-01 01:00:00	55.0	0.6	55.0	55.0	135.0	74.0	55.0	55.0	55.0	7.0	55.0	55.0	28079004
1	2013-11-01 01:00:00	1.5	0.5	1.3	55.0	71.0	83.0	2.0	23.0	16.0	12.0	55.0	8.3	28079008
2	2013-11-01 01:00:00	3.9	55.0	2.8	55.0	49.0	70.0	55.0	55.0	55.0	55.0	55.0	9.0	28079011
3	2013-11-01 01:00:00	55.0	0.5	55.0	55.0	82.0	87.0	3.0	55.0	55.0	55.0	55.0	55.0	28079016
4	2013-11-01 01:00:00	55.0	55.0	55.0	55.0	242.0	111.0	2.0	55.0	55.0	12.0	55.0	55.0	28079017
...
209875	2013-03-01 00:00:00	55.0	0.4	55.0	55.0	8.0	39.0	52.0	55.0	55.0	55.0	55.0	55.0	28079056
209876	2013-03-01 00:00:00	55.0	0.4	55.0	55.0	1.0	11.0	55.0	6.0	55.0	2.0	55.0	55.0	28079057
209877	2013-03-01 00:00:00	55.0	55.0	55.0	55.0	2.0	4.0	75.0	55.0	55.0	55.0	55.0	55.0	28079058
209878	2013-03-01 00:00:00	55.0	55.0	55.0	55.0	2.0	11.0	52.0	55.0	55.0	55.0	55.0	55.0	28079059
209879	2013-03-01 00:00:00	55.0	55.0	55.0	55.0	1.0	10.0	75.0	3.0	55.0	55.0	55.0	55.0	28079060

209880 rows × 14 columns

```
In [611]: b.columns
```

Out[611]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

```
In [612]: c=b.head(20)
c
```

Out[612]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2013-11-01 01:00:00	55.0	0.6	55.0	55.00	135.0	74.0	55.0	55.0	55.0	7.0	55.00	55.0	28079004
1	2013-11-01 01:00:00	1.5	0.5	1.3	55.00	71.0	83.0	2.0	23.0	16.0	12.0	55.00	8.3	28079008
2	2013-11-01 01:00:00	3.9	55.0	2.8	55.00	49.0	70.0	55.0	55.0	55.0	55.0	55.00	9.0	28079011
3	2013-11-01 01:00:00	55.0	0.5	55.0	55.00	82.0	87.0	3.0	55.0	55.0	55.0	55.00	55.0	28079016
4	2013-11-01 01:00:00	55.0	55.0	55.0	55.00	242.0	111.0	2.0	55.0	55.0	12.0	55.00	55.0	28079017
5	2013-11-01 01:00:00	1.0	0.6	0.8	55.00	70.0	70.0	2.0	24.0	55.0	6.0	55.00	5.2	28079018
6	2013-11-01 01:00:00	55.0	0.4	55.0	0.29	51.0	80.0	5.0	23.0	14.0	4.0	1.44	55.0	28079024
7	2013-11-01 01:00:00	55.0	55.0	55.0	0.23	29.0	60.0	4.0	55.0	55.0	55.0	1.51	55.0	28079027
8	2013-11-01 01:00:00	55.0	1.0	55.0	55.00	165.0	107.0	2.0	55.0	55.0	11.0	55.00	55.0	28079035
9	2013-11-01 01:00:00	55.0	0.6	55.0	55.00	63.0	93.0	55.0	11.0	55.0	8.0	55.00	55.0	28079036
10	2013-11-01 01:00:00	1.4	55.0	1.4	55.00	68.0	84.0	55.0	26.0	11.0	6.0	55.00	7.4	28079038
11	2013-11-01 01:00:00	55.0	0.6	55.0	55.00	60.0	82.0	5.0	55.0	55.0	55.0	55.00	55.0	28079039
12	2013-11-01 01:00:00	55.0	55.0	55.0	55.00	69.0	96.0	55.0	29.0	55.0	5.0	55.00	55.0	28079040
13	2013-11-01 01:00:00	55.0	55.0	55.0	55.00	122.0	72.0	55.0	32.0	20.0	55.0	55.00	55.0	28079047
14	2013-11-01 01:00:00	55.0	55.0	55.0	55.00	43.0	81.0	55.0	13.0	10.0	55.0	55.00	55.0	28079048
15	2013-11-01 01:00:00	55.0	55.0	55.0	55.00	132.0	77.0	2.0	55.0	55.0	55.0	55.00	55.0	28079049
16	2013-11-01 01:00:00	55.0	55.0	55.0	55.00	102.0	73.0	55.0	19.0	9.0	55.0	55.00	55.0	28079050
17	2013-11-01 01:00:00	55.0	55.0	55.0	55.00	169.0	74.0	3.0	55.0	55.0	55.0	55.00	55.0	28079054
18	2013-11-01 01:00:00	1.6	55.0	1.4	0.22	62.0	72.0	55.0	20.0	55.0	55.0	1.67	5.4	28079055
19	2013-11-01 01:00:00	55.0	0.8	55.0	55.00	115.0	87.0	5.0	55.0	55.0	55.0	55.00	55.0	28079056


```
In [613]: d=c[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
               'PM10', 'SO_2', 'TCH', 'TOL', 'station']]
d
```

Out[613]:

	BEN	CO	EBE	NMHC	NO_2	O_3	PM10	SO_2	TCH	TOL	station
0	55.0	0.6	55.0	55.00	74.0	55.0	55.0	7.0	55.00	55.0	28079004
1	1.5	0.5	1.3	55.00	83.0	2.0	23.0	12.0	55.00	8.3	28079008
2	3.9	55.0	2.8	55.00	70.0	55.0	55.0	55.0	55.00	9.0	28079011
3	55.0	0.5	55.0	55.00	87.0	3.0	55.0	55.0	55.00	55.0	28079016
4	55.0	55.0	55.0	55.00	111.0	2.0	55.0	12.0	55.00	55.0	28079017
5	1.0	0.6	0.8	55.00	70.0	2.0	24.0	6.0	55.00	5.2	28079018
6	55.0	0.4	55.0	0.29	80.0	5.0	23.0	4.0	1.44	55.0	28079024
7	55.0	55.0	55.0	0.23	60.0	4.0	55.0	55.0	1.51	55.0	28079027
8	55.0	1.0	55.0	55.00	107.0	2.0	55.0	11.0	55.00	55.0	28079035
9	55.0	0.6	55.0	55.00	93.0	55.0	11.0	8.0	55.00	55.0	28079036
10	1.4	55.0	1.4	55.00	84.0	55.0	26.0	6.0	55.00	7.4	28079038
11	55.0	0.6	55.0	55.00	82.0	5.0	55.0	55.0	55.00	55.0	28079039
12	55.0	55.0	55.0	55.00	96.0	55.0	29.0	5.0	55.00	55.0	28079040
13	55.0	55.0	55.0	55.00	72.0	55.0	32.0	55.0	55.00	55.0	28079047
14	55.0	55.0	55.0	55.00	81.0	55.0	13.0	55.0	55.00	55.0	28079048
15	55.0	55.0	55.0	55.00	77.0	2.0	55.0	55.0	55.00	55.0	28079049
16	55.0	55.0	55.0	55.00	73.0	55.0	19.0	55.0	55.00	55.0	28079050
17	55.0	55.0	55.0	55.00	74.0	3.0	55.0	55.0	55.00	55.0	28079054
18	1.6	55.0	1.4	0.22	72.0	55.0	20.0	55.0	1.67	5.4	28079055
19	55.0	0.8	55.0	55.00	87.0	5.0	55.0	55.0	55.00	55.0	28079056

In [614]: `sns.pairplot(d)`


Out[614]: <seaborn.axisgrid.PairGrid at 0x1b6f8031670>



```
In [615]: sns.distplot(a['NO_2'])
```


C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[615]: <AxesSubplot:xlabel='NO_2', ylabel='Density'>
```



```
In [616]: sns.heatmap(d.corr())
```

```
Out[616]: <AxesSubplot:>
```



```
In [617]: x=d[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2']]
y=d['TCH']
```

```
In [618]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [619]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[619]: LinearRegression()
```

In [620]: `print(lr.intercept_)`

-7.686001097830285


In [621]: `coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])`
`coeff`

Out[621]:

	Co-efficient
BEN	43.983082
CO	0.000888
EBE	-43.821223
NMHC	0.977052
NO_2	0.000138

In [622]: `prediction=lr.predict(x_test)`
`plt.scatter(y_test,prediction)`

Out[622]: <matplotlib.collections.PathCollection at 0x1b71ec21bb0>



In [623]: `print(lr.score(x_test,y_test))`

0.0

In [624]: `from sklearn.linear_model import Ridge,Lasso`

In [625]: `rr=Ridge(alpha=10)`
`rr.fit(x_train,y_train)`

Out[625]: Ridge(alpha=10)

In [626]: `rr.score(x_test,y_test)`

Out[626]: 0.0


```
In [627]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[627]: Lasso(alpha=10)
```

```
In [628]: la.score(x_test,y_test)
```

```
Out[628]: 0.0
```

```
In [629]: a1=b.head(7000)
a1
```

```
Out[629]:
```

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2013-11-01 01:00:00	55.0	0.6	55.0	55.0	135.0	74.0	55.0	55.0	55.0	7.0	55.0	55.0	28079004
1	2013-11-01 01:00:00	1.5	0.5	1.3	55.0	71.0	83.0	2.0	23.0	16.0	12.0	55.0	8.3	28079008
2	2013-11-01 01:00:00	3.9	55.0	2.8	55.0	49.0	70.0	55.0	55.0	55.0	55.0	55.0	9.0	28079011
3	2013-11-01 01:00:00	55.0	0.5	55.0	55.0	82.0	87.0	3.0	55.0	55.0	55.0	55.0	55.0	28079016
4	2013-11-01 01:00:00	55.0	55.0	55.0	55.0	242.0	111.0	2.0	55.0	55.0	12.0	55.0	55.0	28079017
...
6995	2013-11-13 04:00:00	55.0	0.2	55.0	55.0	1.0	8.0	40.0	55.0	55.0	55.0	55.0	55.0	28079039
6996	2013-11-13 04:00:00	55.0	55.0	55.0	55.0	1.0	5.0	55.0	3.0	55.0	1.0	55.0	55.0	28079040
6997	2013-11-13 04:00:00	55.0	55.0	55.0	55.0	1.0	6.0	55.0	3.0	2.0	55.0	55.0	55.0	28079047
6998	2013-11-13 04:00:00	55.0	55.0	55.0	55.0	1.0	9.0	55.0	5.0	1.0	55.0	55.0	55.0	28079048
6999	2013-11-13 04:00:00	55.0	55.0	55.0	55.0	1.0	9.0	43.0	55.0	55.0	55.0	55.0	55.0	28079049

7000 rows × 14 columns

```
In [630]: e=a1[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
'PM10', 'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [631]: f=e.iloc[:,0:14]
g=e.iloc[:, -1]
```

```
In [632]: h=StandardScaler().fit_transform(f)
```

```
In [633]: logr=LogisticRegression(max_iter=10000)
logr.fit(h,g)
```

```
Out[633]: LogisticRegression(max_iter=10000)
```

```
In [634]: from sklearn.model_selection import train_test_split
h_train,h_test,g_train,g_test=train_test_split(h,g,test_size=0.3)
```

```
In [635]: i=[[10,20,30,40,50,60,15,26,37,47,58]]
```

```
In [636]: prediction=logr.predict(i)
print(prediction)
```

```
[28079059]
```

```
In [637]: logr.classes_
```

```
Out[637]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,
                28079024, 28079027, 28079035, 28079036, 28079038, 28079039,
                28079040, 28079047, 28079048, 28079049, 28079050, 28079054,
                28079055, 28079056, 28079057, 28079058, 28079059, 28079060],
              dtype=int64)
```

```
In [638]: logr.predict_proba(i)[0][0]
```

```
Out[638]: 0.0
```

```
In [639]: logr.predict_proba(i)[0][1]
```

```
Out[639]: 0.0
```

```
In [640]: logr.score(h_test,g_test)
```

```
Out[640]: 0.9485714285714286
```

```
In [641]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[641]: ElasticNet()
```

```
In [642]: print(en.coef_)
```

```
[-0.          0.         -0.          0.97442165  0.          ]
```

```
In [643]: print(en.intercept_)
```

```
1.3838448105340149
```

```
In [644]: prediction=en.predict(x_test)
print(en.score(x_test,y_test))
```

```
0.0
```

```
In [645]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(h_train,g_train)
```

Out[645]: RandomForestClassifier()

```
In [646]: parameters={'max_depth':[1,2,3,4,5],
'min_samples_leaf':[5,10,15,20,25],
'n_estimators':[10,20,30,40,50]
}
```

```
In [647]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(h_train,g_train)
```

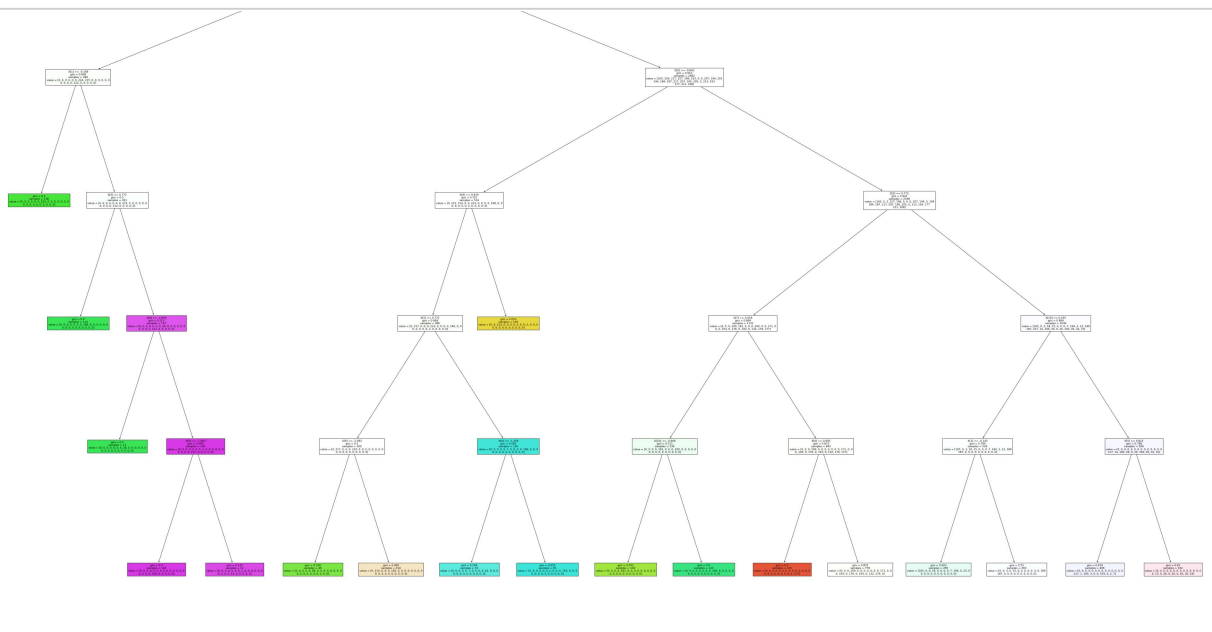
Out[647]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')

```
In [648]: grid_search.best_score_
```

Out[648]: 0.9977551020408163

```
In [649]: rfc_best=grid_search.best_estimator_
```

```
In [651]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,50))
plot_tree(rfc_best.estimators_[2],filled=True)
```



Conclusion: from this data set i observed that the ridge has the highest accuracy of 0.9961224489795919

In []: