

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import re
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
```

```
In [2]: a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\madrid_2016\
a
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2016-11-01 01:00:00	NaN	0.7	NaN	NaN	153.0	77.0	NaN	NaN	NaN	7.0	NaN	NaN	28079004
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	28079008
2	2016-11-01 01:00:00	5.9	NaN	7.5	NaN	297.0	139.0	NaN	NaN	NaN	NaN	NaN	26.0	28079011
3	2016-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	113.0	2.0	NaN	NaN	NaN	NaN	NaN	28079016
4	2016-11-01 01:00:00	NaN	NaN	NaN	NaN	275.0	127.0	2.0	NaN	NaN	18.0	NaN	NaN	28079017
...
209491	2016-07-01 00:00:00	NaN	0.2	NaN	NaN	2.0	29.0	73.0	NaN	NaN	NaN	NaN	NaN	28079056
209492	2016-07-01 00:00:00	NaN	0.3	NaN	NaN	1.0	29.0	NaN	36.0	NaN	5.0	NaN	NaN	28079057
209493	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	1.0	19.0	71.0	NaN	NaN	NaN	NaN	NaN	28079058
209494	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	6.0	17.0	85.0	NaN	NaN	NaN	NaN	NaN	28079059
209495	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	2.0	46.0	61.0	34.0	NaN	NaN	NaN	NaN	28079060

209496 rows × 14 columns

```
In [3]: a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209496 entries, 0 to 209495
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        209496 non-null object  
 1   BEN         50755 non-null  float64
 2   CO          85999 non-null  float64
 3   EBE         50335 non-null  float64
 4   NMHC        25970 non-null  float64
 5   NO          208614 non-null float64
 6   NO_2        208614 non-null float64
 7   O_3         121197 non-null float64
 8   PM10        102892 non-null float64
 9   PM25        52165 non-null  float64
10   SO_2        86023 non-null  float64
11   TCH         25970 non-null  float64
12   TOL         50662 non-null  float64
13   station     209496 non-null int64  
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [4]: b=a.fillna(value=86)
b
```

Out[4]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2016-11-01 01:00:00	86.0	0.7	86.0	86.00	153.0	77.0	86.0	86.0	86.0	7.0	86.00	86.0	28079004
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	28079008
2	2016-11-01 01:00:00	5.9	86.0	7.5	86.00	297.0	139.0	86.0	86.0	86.0	86.00	26.0		28079011
3	2016-11-01 01:00:00	86.0	1.0	86.0	86.00	154.0	113.0	2.0	86.0	86.0	86.0	86.00	86.0	28079016
4	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	275.0	127.0	2.0	86.0	86.0	18.0	86.00	86.0	28079017
...
209491	2016-07-01 00:00:00	86.0	0.2	86.0	86.00	2.0	29.0	73.0	86.0	86.0	86.0	86.00	86.0	28079056
209492	2016-07-01 00:00:00	86.0	0.3	86.0	86.00	1.0	29.0	86.0	36.0	86.0	5.0	86.00	86.0	28079057
209493	2016-07-01 00:00:00	86.0	86.0	86.0	86.00	1.0	19.0	71.0	86.0	86.0	86.0	86.00	86.0	28079058
209494	2016-07-01 00:00:00	86.0	86.0	86.0	86.00	6.0	17.0	85.0	86.0	86.0	86.0	86.00	86.0	28079059
209495	2016-07-01 00:00:00	86.0	86.0	86.0	86.00	2.0	46.0	61.0	34.0	86.0	86.0	86.00	86.0	28079060

209496 rows × 14 columns

```
In [5]: b.columns
```

Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

```
In [6]: c=b.head(30)  
c
```

Out[6]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2016-11-01 01:00:00	86.0	0.7	86.0	86.00	153.0	77.0	86.0	86.0	86.0	7.0	86.00	86.000000	28079004
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.400000	28079008
2	2016-11-01 01:00:00	5.9	86.0	7.5	86.00	297.0	139.0	86.0	86.0	86.0	86.0	86.00	26.000000	28079011
3	2016-11-01 01:00:00	86.0	1.0	86.0	86.00	154.0	113.0	2.0	86.0	86.0	86.0	86.00	86.000000	28079016
4	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	275.0	127.0	2.0	86.0	86.0	18.0	86.00	86.000000	28079017
5	2016-11-01 01:00:00	0.9	0.5	0.5	86.00	66.0	82.0	1.0	27.0	86.0	8.0	86.00	6.000000	28079018
6	2016-11-01 01:00:00	0.7	0.8	0.4	0.13	57.0	66.0	3.0	23.0	15.0	4.0	1.35	5.000000	28079024
7	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	52.0	78.0	1.0	86.0	86.0	86.0	86.00	86.000000	28079027
8	2016-11-01 01:00:00	86.0	1.2	86.0	86.00	205.0	85.0	6.0	86.0	86.0	86.0	86.00	86.000000	28079035
9	2016-11-01 01:00:00	86.0	0.7	86.0	86.00	114.0	91.0	86.0	37.0	86.0	6.0	86.00	86.000000	28079036
10	2016-11-01 01:00:00	2.5	86.0	3.3	86.00	166.0	114.0	86.0	45.0	27.0	8.0	86.00	16.299999	28079038
11	2016-11-01 01:00:00	86.0	2.4	86.0	86.00	475.0	165.0	5.0	86.0	86.0	86.0	86.00	86.000000	28079039
12	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	74.0	109.0	86.0	38.0	86.0	8.0	86.00	86.000000	28079040
13	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	168.0	93.0	86.0	41.0	26.0	86.0	86.00	86.000000	28079047
14	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	150.0	85.0	86.0	31.0	21.0	86.0	86.00	86.000000	28079048
15	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	81.0	97.0	1.0	86.0	86.0	86.0	86.00	86.000000	28079049
16	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	307.0	134.0	86.0	55.0	35.0	86.0	86.00	86.000000	28079050
17	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	160.0	113.0	1.0	86.0	86.0	86.0	86.00	86.000000	28079054

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
18	2016-11-01 01:00:00	1.4	86.0	1.3	0.20	72.0	84.0	86.0	21.0	86.0	86.0	1.50	6.900000	28079055
19	2016-11-01 01:00:00	86.0	0.5	86.0	86.00	67.0	83.0	6.0	86.0	86.0	86.0	86.00	86.000000	28079056
20	2016-11-01 01:00:00	86.0	1.0	86.0	86.00	181.0	109.0	86.0	42.0	86.0	13.0	86.00	86.000000	28079057
21	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	4.0	36.0	11.0	86.0	86.0	86.0	86.00	86.000000	28079058
22	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	99.0	66.0	7.0	86.0	86.0	86.0	86.00	86.000000	28079059
23	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	79.0	103.0	4.0	29.0	86.0	86.0	86.00	86.000000	28079060
24	2016-11-01 02:00:00	86.0	0.6	86.0	86.00	116.0	65.0	86.0	86.0	86.0	6.0	86.00	86.000000	28079004
25	2016-11-01 02:00:00	2.7	1.0	2.1	0.40	139.0	114.0	4.0	37.0	21.0	14.0	2.30	15.000000	28079008
26	2016-11-01 02:00:00	4.7	86.0	5.6	86.00	111.0	97.0	86.0	86.0	86.0	86.0	86.00	16.700001	28079011
27	2016-11-01 02:00:00	86.0	0.7	86.0	86.00	67.0	90.0	2.0	86.0	86.0	86.0	86.00	86.000000	28079016
28	2016-11-01 02:00:00	86.0	86.0	86.0	86.00	99.0	84.0	2.0	86.0	86.0	13.0	86.00	86.000000	28079017
29	2016-11-01 02:00:00	0.5	0.5	0.2	86.00	61.0	73.0	1.0	17.0	86.0	7.0	86.00	3.300000	28079018


```
In [7]: d=c[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
            'PM10', 'SO_2', 'TCH', 'TOL', 'station']]
d
```

Out[7]:

	BEN	CO	EBE	NMHC	NO_2	O_3	PM10	SO_2	TCH	TOL	station
0	86.0	0.7	86.0	86.00	77.0	86.0	86.0	7.0	86.00	86.000000	28079004
1	3.1	1.1	2.0	0.53	144.0	4.0	46.0	18.0	2.44	14.400000	28079008
2	5.9	86.0	7.5	86.00	139.0	86.0	86.0	86.0	86.00	26.000000	28079011
3	86.0	1.0	86.0	86.00	113.0	2.0	86.0	86.0	86.00	86.000000	28079016
4	86.0	86.0	86.0	86.00	127.0	2.0	86.0	18.0	86.00	86.000000	28079017
5	0.9	0.5	0.5	86.00	82.0	1.0	27.0	8.0	86.00	6.000000	28079018
6	0.7	0.8	0.4	0.13	66.0	3.0	23.0	4.0	1.35	5.000000	28079024
7	86.0	86.0	86.0	86.00	78.0	1.0	86.0	86.0	86.00	86.000000	28079027
8	86.0	1.2	86.0	86.00	85.0	6.0	86.0	86.0	86.00	86.000000	28079035
9	86.0	0.7	86.0	86.00	91.0	86.0	37.0	6.0	86.00	86.000000	28079036
10	2.5	86.0	3.3	86.00	114.0	86.0	45.0	8.0	86.00	16.299999	28079038
11	86.0	2.4	86.0	86.00	165.0	5.0	86.0	86.0	86.00	86.000000	28079039
12	86.0	86.0	86.0	86.00	109.0	86.0	38.0	8.0	86.00	86.000000	28079040
13	86.0	86.0	86.0	86.00	93.0	86.0	41.0	86.0	86.00	86.000000	28079047
14	86.0	86.0	86.0	86.00	85.0	86.0	31.0	86.0	86.00	86.000000	28079048
15	86.0	86.0	86.0	86.00	97.0	1.0	86.0	86.0	86.00	86.000000	28079049
16	86.0	86.0	86.0	86.00	134.0	86.0	55.0	86.0	86.00	86.000000	28079050
17	86.0	86.0	86.0	86.00	113.0	1.0	86.0	86.0	86.00	86.000000	28079054
18	1.4	86.0	1.3	0.20	84.0	86.0	21.0	86.0	1.50	6.900000	28079055
19	86.0	0.5	86.0	86.00	83.0	6.0	86.0	86.0	86.00	86.000000	28079056
20	86.0	1.0	86.0	86.00	109.0	86.0	42.0	13.0	86.00	86.000000	28079057
21	86.0	86.0	86.0	86.00	36.0	11.0	86.0	86.0	86.00	86.000000	28079058
22	86.0	86.0	86.0	86.00	66.0	7.0	86.0	86.0	86.00	86.000000	28079059
23	86.0	86.0	86.0	86.00	103.0	4.0	29.0	86.0	86.00	86.000000	28079060
24	86.0	0.6	86.0	86.00	65.0	86.0	86.0	6.0	86.00	86.000000	28079004
25	2.7	1.0	2.1	0.40	114.0	4.0	37.0	14.0	2.30	15.000000	28079008
26	4.7	86.0	5.6	86.00	97.0	86.0	86.0	86.0	86.00	16.700001	28079011
27	86.0	0.7	86.0	86.00	90.0	2.0	86.0	86.0	86.00	86.000000	28079016
28	86.0	86.0	86.0	86.00	84.0	2.0	86.0	13.0	86.00	86.000000	28079017
29	0.5	0.5	0.2	86.00	73.0	1.0	17.0	7.0	86.00	3.300000	28079018

```
In [8]: sns.pairplot(d)
```


```
Out[8]: <seaborn.axisgrid.PairGrid at 0x1d524e13b50>
```




```
In [9]: sns.distplot(a['NO_2'])
```


C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[9]: <AxesSubplot:xlabel='NO_2', ylabel='Density'>
```



```
In [10]: sns.heatmap(d.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: x=d[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2']]
y=d['TCH']
```

```
In [12]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [13]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[13]: LinearRegression()
```

In [14]: `print(lr.intercept_)`

1.4748792800566548


In [15]: `coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])`
`coeff`

Out[15]:

	Co-efficient
BEN	0.134493
CO	-0.000405
EBE	-0.134567
NMHC	0.981957
NO_2	0.001019

In [16]: `prediction=lr.predict(x_test)`
`plt.scatter(y_test,prediction)`

Out[16]: <matplotlib.collections.PathCollection at 0xd52d895700>



In [17]: `print(lr.score(x_test,y_test))`

0.9999806673343318

In [18]: `from sklearn.linear_model import Ridge,Lasso`

In [19]: `rr=Ridge(alpha=10)`
`rr.fit(x_train,y_train)`

Out[19]: Ridge(alpha=10)

In [20]: `rr.score(x_test,y_test)`

Out[20]: 0.9999768962712836

```
In [21]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[21]: Lasso(alpha=10)
```

```
In [22]: la.score(x_test,y_test)
```

```
Out[22]: 0.999678913074025
```

```
In [23]: a1=b.head(7000)
a1
```

```
Out[23]:
```

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2016-11-01 01:00:00	86.0	0.7	86.0	86.00	153.0	77.0	86.0	86.0	86.0	7.0	86.00	86.0	28079004
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	28079008
2	2016-11-01 01:00:00	5.9	86.0	7.5	86.00	297.0	139.0	86.0	86.0	86.0	86.00	86.00	26.0	28079011
3	2016-11-01 01:00:00	86.0	1.0	86.0	86.00	154.0	113.0	2.0	86.0	86.0	86.0	86.00	86.0	28079016
4	2016-11-01 01:00:00	86.0	86.0	86.0	86.00	275.0	127.0	2.0	86.0	86.0	18.0	86.00	86.0	28079017
...
6995	2016-11-13 04:00:00	86.0	0.7	86.0	86.00	96.0	71.0	5.0	86.0	86.0	86.0	86.00	86.0	28079039
6996	2016-11-13 04:00:00	86.0	86.0	86.0	86.00	45.0	70.0	86.0	26.0	86.0	9.0	86.00	86.0	28079040
6997	2016-11-13 04:00:00	86.0	86.0	86.0	86.00	87.0	70.0	86.0	28.0	23.0	86.0	86.00	86.0	28079047
6998	2016-11-13 04:00:00	86.0	86.0	86.0	86.00	66.0	59.0	86.0	33.0	26.0	86.0	86.00	86.0	28079048
6999	2016-11-13 04:00:00	86.0	86.0	86.0	86.00	98.0	53.0	1.0	86.0	86.0	86.0	86.00	86.0	28079049

7000 rows × 14 columns

```
In [24]: e=a1[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
'PM10', 'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [25]: f=e.iloc[:,0:14]
g=e.iloc[:, -1]
```

```
In [26]: h=StandardScaler().fit_transform(f)
```

```
In [27]: logr=LogisticRegression(max_iter=10000)
logr.fit(h,g)
```

```
Out[27]: LogisticRegression(max_iter=10000)
```

```
In [28]: from sklearn.model_selection import train_test_split
h_train,h_test,g_train,g_test=train_test_split(h,g,test_size=0.3)
```

```
In [29]: i=[[10,20,30,40,50,60,15,26,37,47,58]]
```

```
In [30]: prediction=logr.predict(i)
print(prediction)

[28079059]
```

```
In [31]: logr.classes_
```

```
Out[31]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,
                28079024, 28079027, 28079035, 28079036, 28079038, 28079039,
                28079040, 28079047, 28079048, 28079049, 28079050, 28079054,
                28079055, 28079056, 28079057, 28079058, 28079059, 28079060],
              dtype=int64)
```

```
In [32]: logr.predict_proba(i)[0][0]
```

```
Out[32]: 0.0
```

```
In [33]: logr.predict_proba(i)[0][1]
```

```
Out[33]: 0.0
```

```
In [34]: logr.score(h_test,g_test)
```

```
Out[34]: 0.9452380952380952
```

```
In [35]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[35]: ElasticNet()
```

```
In [36]: print(en.coef_)
```

```
[ 0.00000000e+00 -4.53645120e-04  5.21275022e-04  9.79350801e-01
 0.00000000e+00]
```

```
In [37]: print(en.intercept_)
```

```
1.7459678449370415
```

```
In [38]: prediction=en.predict(x_test)
print(en.score(x_test,y_test))
```

0.9999713903612383

```
In [39]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(h_train,g_train)
```

Out[39]: RandomForestClassifier()

```
In [40]: parameters={'max_depth':[1,2,3,4,5],
'min_samples_leaf':[5,10,15,20,25],
'n_estimators':[10,20,30,40,50]
}
```

```
In [41]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(h_train,g_train)
```

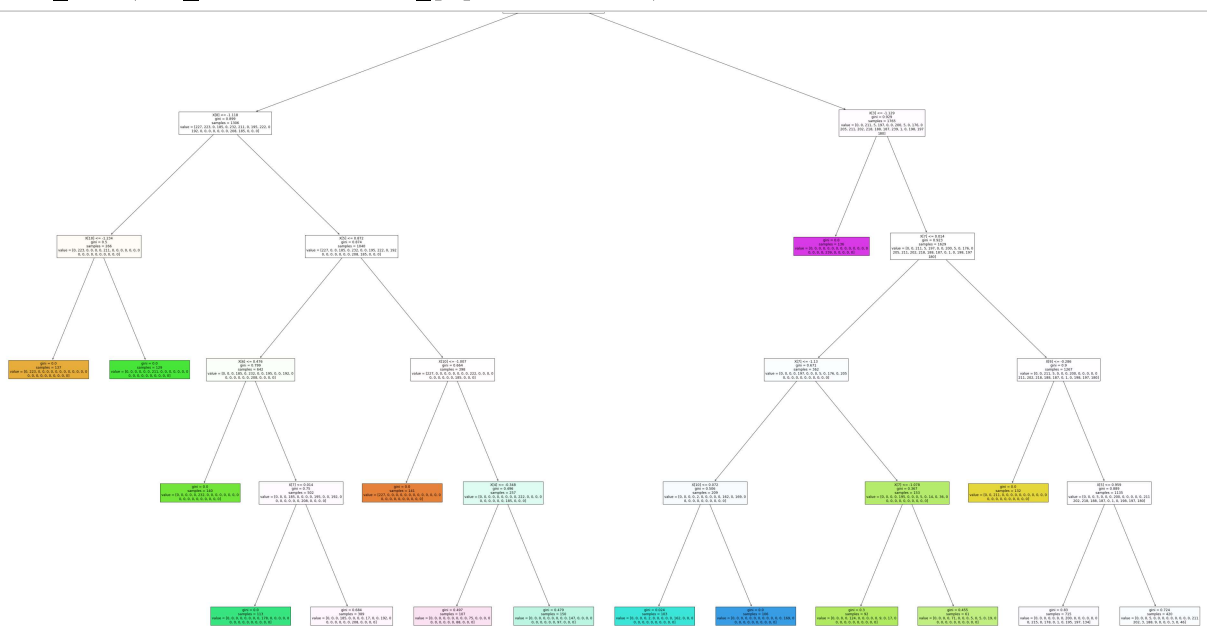
Out[41]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')

```
In [42]: grid_search.best_score_
```

Out[42]: 0.9936734693877551

```
In [43]: rfc_best=grid_search.best_estimator_
```

```
In [44]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,50))
plot_tree(rfc_best.estimators_[2],filled=True)
```



Conclusion: from this data set i observed that the ridge has the highest accuracy of 0.9999768962712836

In []: